

```
In [46]: import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import numpy as np
```

```
In [47]: data = pd.read_csv('lifesat.csv')
```

```
In [48]: # Q1 (5 marks) Load the dataset lifesat.csv and display the first 5  
data.head(5)
```

```
Out[48]:    Country  GDP per capita (USD)  Life satisfaction  
0    Russia        26456.387938          5.8  
1    Greece        27287.083401          5.4  
2    Turkey        28384.987785          5.5  
3    Latvia        29932.493910          5.9  
4    Hungary       31007.768407          5.6
```

```
In [49]: # Q2 (5 marks) Print basic info and summary statistics .  
data.info()
```

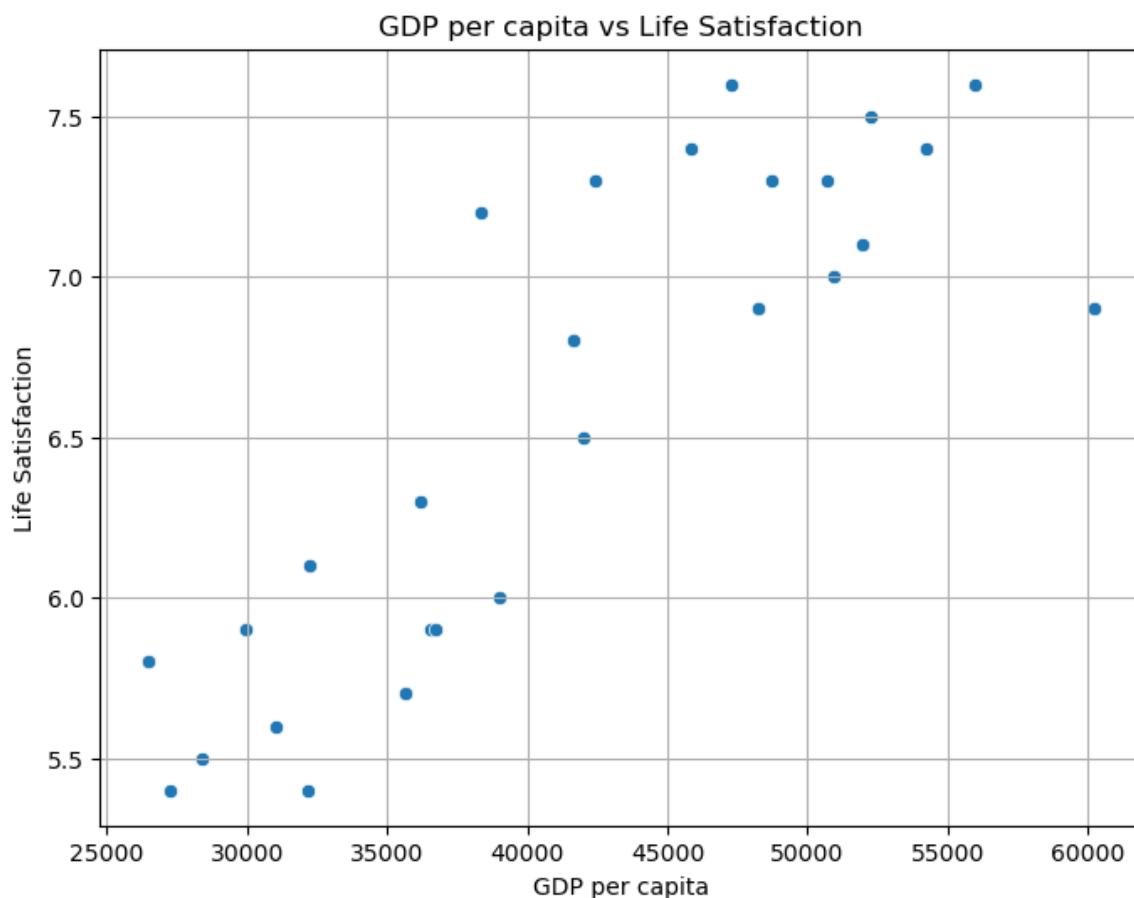
```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 27 entries, 0 to 26  
Data columns (total 3 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   Country          27 non-null      object    
 1   GDP per capita (USD) 27 non-null      float64  
 2   Life satisfaction  27 non-null      float64  
dtypes: float64(2), object(1)  
memory usage: 780.0+ bytes
```

```
In [50]: data.describe()
```

```
Out[50]:    GDP per capita (USD)  Life satisfaction  
count      27.000000          27.000000  
mean      41564.521771         6.566667  
std       9631.452319         0.765607  
min      26456.387938         5.400000  
25%      33938.289305         5.900000  
50%      41627.129269         6.800000  
75%      49690.580269         7.300000  
max      60235.728492         7.600000
```

In [51]:

```
# Q3 (10 marks) Display a scatter plot for GDP per capita vs Life Satisfaction
plt.figure(figsize=(8,6))
sns.scatterplot(data=data, x = 'GDP per capita (USD)', y = 'Life satisfaction')
plt.xlabel('GDP per capita')
plt.ylabel('Life Satisfaction')
plt.title('GDP per capita vs Life Satisfaction')
plt.grid(True)
plt.show()
```



Discuss the observed relationship:

In the plot, we can see that GDP per capita is an important factor influencing life satisfaction. However, GDP alone is not enough to fully explain life satisfaction, because in some countries other factors, such as culture, healthcare, and social support, can play a more significant role.

In [52]:

```
# Q4 (5 marks) Extract input (X) and target (y). Print their shape
X = data[['GDP per capita (USD)']] # [[[]]] -> DataFrame
y = data['Life satisfaction']

print("Shape of X:", X.shape)
print("Shape of y:", y.shape)
```

Shape of X: (27, 1)

Shape of y: (27,)

In [53]:

```
# Q5 (10 marks) Train a Linear Regression model & Display coefficients
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
model.fit(X, y)

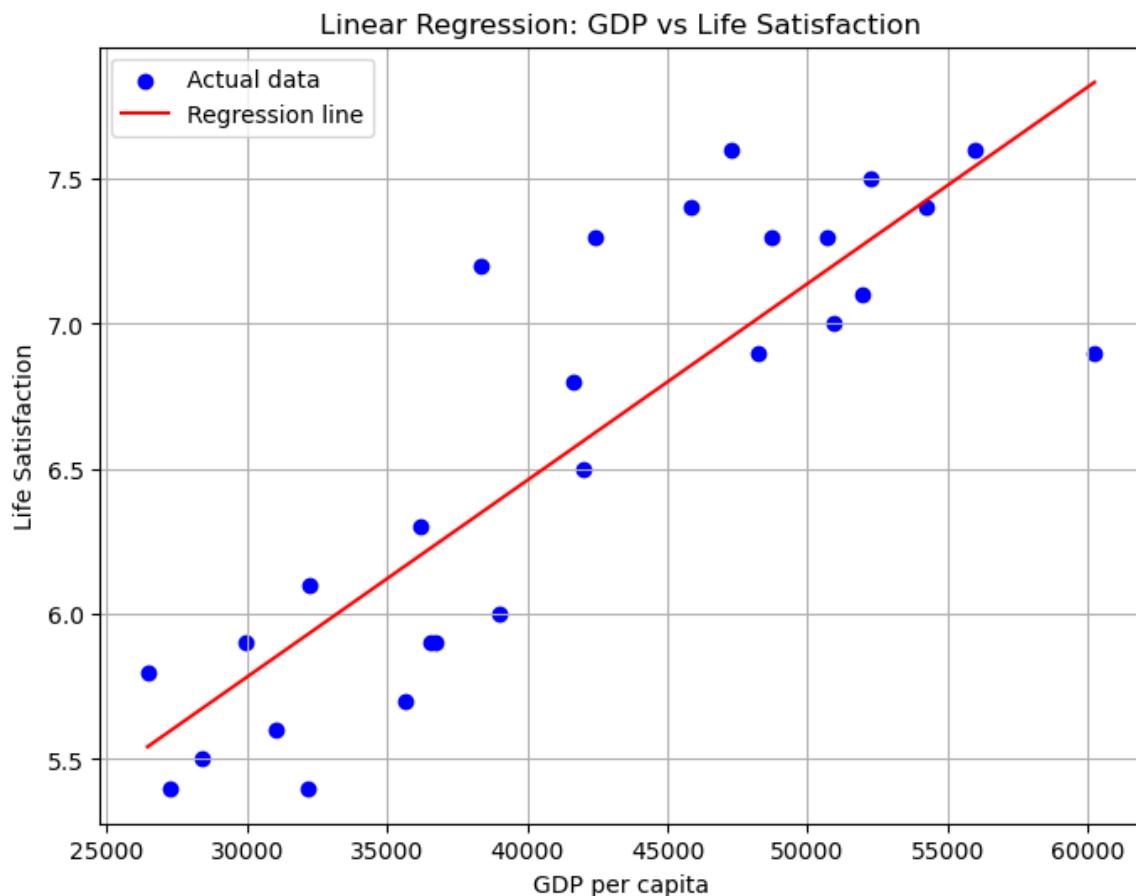
print("Coefficient (slope):", model.coef_[0])
print("Intercept:", model.intercept_)
```

Coefficient (slope): 6.778899694341222e-05
 Intercept: 3.7490494273769093

In [54]: # Q6 (10 marks) Plot the predicted regression line from the model a

```
y_pred = model.predict(X) # prediction

plt.figure(figsize=(8,6))
plt.scatter(X, y, label='Actual data', color='blue')
plt.plot(X, y_pred, color='red', label='Regression line')
plt.xlabel('GDP per capita')
plt.ylabel('Life Satisfaction')
plt.title('Linear Regression: GDP vs Life Satisfaction')
plt.legend()
plt.grid(True)
plt.show()
```



In [55]: # Q7 (5 marks) Predict Life Satisfaction for GDP = 37,655.2 USD. Consider the data above.
 gdp = 37655.2
 life_satisfaction_pred = model.predict([[gdp]])
 print("Predicted Life Satisfaction for GDP = 37,655.2 USD:", life_s

Predicted Life Satisfaction for GDP = 37,655.2 USD: 6.301657665080485

```
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
warnings.warn(
```

In [56]: # Q8 (5 marks) Train a KNeighborsRegressor (n_neighbors=3).
from sklearn.neighbors import KNeighborsRegressor

knn_model = KNeighborsRegressor(n_neighbors=3)
knn_model.fit(X, y)

Out[56]: ▾ KNeighborsRegressor ① ⓘ

► Parameters

In [57]: # Q9 (10 marks) Predict Life Satisfaction for GDP = 37,655.2 USD and
gdp_value = 37655.2
life_satisfaction_knn = knn_model.predict([[gdp_value]])
print("Predicted Life Satisfaction (KNN) for GDP = 37,655.2 USD:",

print()

life_satisfaction_lr = model.predict(pd.DataFrame({'GDP per capita':
print("Predicted Life Satisfaction (Linear Regression) for GDP = 37

Predicted Life Satisfaction (KNN) for GDP = 37,655.2 USD: 6.33333333
3333333

Predicted Life Satisfaction (Linear Regression) for GDP = 37,655.2 USD: 6.301657665080485

```
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
warnings.warn(
```

In [58]: # Q10 (10 marks)
Use n_neighbors 1, 3, 5, and 10 and print the predicted values of
Plot the results using a line plot.
gdp_value = 37655.2
neighbors_list = [1, 3, 5, 10]
predictions = {}

for k in neighbors_list:
 knn = KNeighborsRegressor(n_neighbors=k)
 knn.fit(X, y)
 pred = knn.predict([[gdp_value]])[0]
 predictions[k] = pred
 print(f"Predicted Life Satisfaction (KNN, n_neighbors={k}) = {pred}")

plt.figure(figsize=(8,6))
plt.plot(list(predictions.keys()), list(predictions.values()), marker='o')
plt.xlabel('Number of Neighbors (n_neighbors)')
plt.ylabel('Predicted Life Satisfaction')

```
plt.title('KNN Predictions for Different n_neighbors')
plt.grid(True)
plt.show()
```

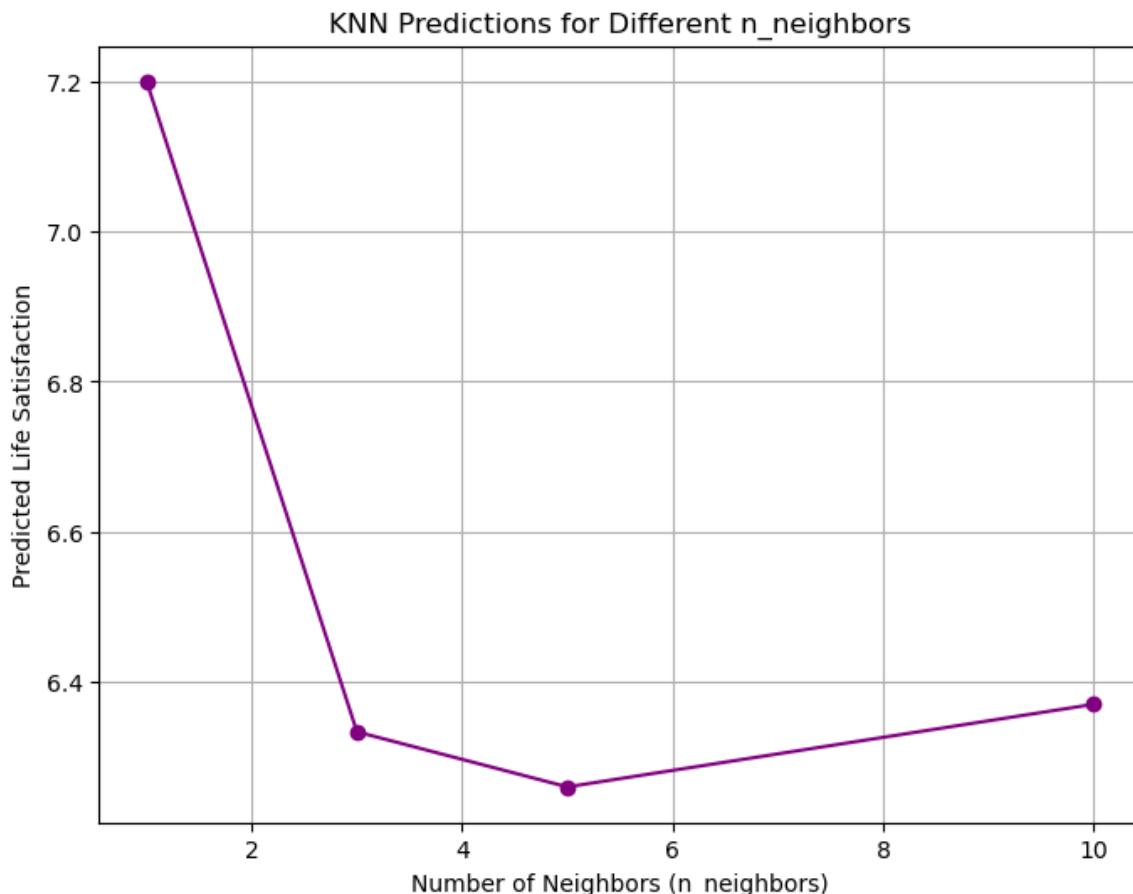
```
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
    warnings.warn(
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
    warnings.warn(
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
    warnings.warn(
/opt/conda/envs/anaconda-2025.12-py312/lib/python3.12/site-packages/
sklearn/utils/validation.py:2749: UserWarning: X does not have valid
feature names, but KNeighborsRegressor was fitted with feature names
    warnings.warn(
```

Predicted Life Satisfaction (KNN, n_neighbors=1) = 7.20

Predicted Life Satisfaction (KNN, n_neighbors=3) = 6.33

Predicted Life Satisfaction (KNN, n_neighbors=5) = 6.26

Predicted Life Satisfaction (KNN, n_neighbors=10) = 6.37



In []:

In []: